

Fox ML Infrastructure — Client Onboarding Guide

Document Hierarchy: This document is provided for guidance only. In case of any conflict, `COMMERCIAL_LICENSE.md` is the authoritative and controlling document for all commercial licensing terms and obligations, and `LICENSE` (AGPL-3.0) controls for open-source licensing.

Commercial Licensees Only: This document applies only to customers with a valid commercial license. AGPL-3.0 users are not eligible for the policies, services, or terms described here.

No Rights Granted: This document does not grant any license rights. All rights and restrictions are defined solely in `COMMERCIAL_LICENSE.md`.

This guide helps commercial licensees get started with Fox ML Infrastructure. It covers setup, configuration, integration, and how to request custom features.

1. Prerequisites

1.1 Required Dependencies

System Requirements: - Linux (Ubuntu 22.04+ recommended) or macOS - Python 3.9+ - Git - Sufficient disk space for datasets and models

Python Dependencies: - Core dependencies are listed in `requirements.txt` (included in repository) - Additional dependencies may be required based on model types used

Hardware Recommendations: - CPU: Multi-core processor (8+ cores recommended) - RAM: 16GB+ (32GB+ for large datasets) - GPU: Optional but recommended for deep learning models (CUDA-compatible)

1.2 Access Setup

Repository Access: - Commercial licensees receive access to private GitHub repositories - Access is granted via GitHub organization or repository-level permissions - See `LEGAL/ENTERPRISE_DELIVERY.md` for repository structure details

Support Access: - Support email: jenn.lewis5789@gmail.com - Support tier determines response times (see `LEGAL/SUPPORT_POLICY.md`)

2. Initial Setup

2.1 Repository Cloning

Enterprise Base Repository:

```
git clone https://github.com/Fox-ML-infrastructure/fox-v1-infra-enterprise.git
cd fox-v1-infra-enterprise
git checkout v1.0.0 # or latest version tag
```

Client-Specific Repository (if applicable):

```
git clone https://github.com/Fox-ML-infrastructure/client-<your-org>-fox-infra.git
cd client-<your-org>-fox-infra
```

2.2 Environment Setup

Python Virtual Environment:

```
python3 -m venv venv
source venv/bin/activate # On Linux/macOS
pip install -r requirements.txt
```

Configuration: - Copy example configuration files to your working directory - Review configuration options in config/ directory - See docs/01_tutorials/configuration/CONFIG_BASICS.md for configuration details

3. Directory Structure

3.1 Expected Structure

Core Directories:

```
fox-v1-infra/
|-- ALPACA_trading/      # Trading integration modules
|-- CONFIG/              # Configuration files
|-- DATA_PROCESSING/    # Data processing pipelines
|-- TRAINING/            # Model training framework
|-- docs/                # Documentation
|-- config/              # Runtime configuration
++-- tests/              # Test suite
```

Client-Specific Structure:

```
client-<your-org>-fox-infra/
|-- config/              # Client-specific configurations
|-- custom/              # Custom modules and features
|-- deployments/         # Deployment recipes
++-- strategies/         # Client-specific strategies (if applicable)
```

3.2 Data Directory

Recommended Data Structure:

```
data/
|-- raw/                 # Raw market data
|-- processed/           # Processed features
|-- models/              # Trained models
++-- results/             # Backtest results and outputs
```

4. Market Data Integration

4.1 Data Sources

Supported Data Sources: - yfinance (default, for equity data) - Custom data sources (via data adapters) - Client-provided data files

4.2 Configuring Data Sources

Configuration Example:

```
data:  
  source: yfinance  
  auto_adjust: false  
  cache: true  
  tickers: [SPY, QQQ, IWM]
```

Custom Data Integration: - Implement data adapter following the existing adapter pattern - See [DATA_PROCESSING/](#) for examples - Contact support for integration assistance

5. Model Configuration

5.1 Available Models

Model Types: - LightGBM, XGBoost (gradient boosting) - MLP, LSTM, Transformer (deep learning) - Ensemble models - Multi-task models - Probabilistic models (NGBoost, Quantile regression)

See [docs/02_reference/models/](#) for complete model documentation.

5.2 Configuring Models

Configuration Example:

```
model:  
  type: lightgbm  
  variant: conservative  
  params:  
    n_estimators: 100  
    learning_rate: 0.05
```

Configuration Files: - Model configs are in [CONFIG/](#) directory - Variants available: `conservative`, `balanced`, `aggressive` - See [docs/01_tutorials/configuration/ADVANCED_CONFIG.md](#) for advanced configuration

6. Running Your First Pipeline

6.1 Basic Workflow

Step 1: Data Processing

```
python -m DATA_PROCESSING.pipelines.process_data \
    --input data/raw \
    --output data/processed \
    --config config/data_config.yaml
```

Step 2: Feature Engineering

```
python -m DATA_PROCESSING.features.build_features \
    --input data/processed \
    --output data/features \
    --config config/features_config.yaml
```

Step 3: Model Training

```
python -m TRAINING.train \
    --data data/features \
    --model lightgbm \
    --output models/ \
    --config config/model_config.yaml
```

6.2 Documentation

Tutorials: - [docs/01_tutorials/pipelines/FIRST_PIPELINE_RUN.md](#) — First pipeline walkthrough - [docs/01_tutorials/training/BASIC_TRAINING.md](#) — Training basics - [docs/00_executive/GETTING_STARTED.md](#) — Quick start guide

7. Custom Features & Integration

7.1 Requesting Custom Features

Process: 1. **Contact support** — Email jenn.lewis5789@gmail.com with feature request 2. **Scoping** — Feature request is evaluated and scoped 3. **Statement of Work (SOW)** — Custom features require a separate SOW (see [LEGAL/consulting/STATEMENT_OF_WORK_TEMPLATE.md](#)) 4. **Development** — Feature is developed in client-specific repository 5. **Delivery** — Feature is delivered via client repository update

See [LEGAL/consulting/CONSULTING_PRICING.md](#) for pricing information.

7.2 Integration with Existing Systems

Integration Support: - **Architecture review** — Enterprise/Premium support includes pre-purchase architectural discussions - **Integration guidance** — Support can provide guidance on integrating with client systems - **Custom adapters** — Custom data adapters or integrations can be developed via SOW

8. Receiving Updates

8.1 Version Updates

Update Process: 1. **Check release notes** — Review CHANGELOG_ENTERPRISE.md for changes 2. **Review migration notes** — Check for breaking changes or migration requirements 3. **Update version tag** — Update your repository to the new version tag 4. **Test in development** — Test updates in a development environment first 5. **Update production** — Update production after successful testing

Example:

```
# In enterprise base repository
git fetch
git checkout v1.1.0 # Update to new version

# In client repository (if applicable)
git pull origin main # Pull any client-specific updates
```

See [LEGAL/RELEASE_POLICY.md](#) for versioning strategy and update recommendations.

8.2 Update Frequency

Recommended Schedule: - **Patch releases** — Update within 1-2 weeks (especially security patches) - **Minor releases** — Update within 1-2 months (test in development first) - **Major releases** — Update when ready (plan migration, test thoroughly)

9. Support & Resources

9.1 Support Channels

Email Support: - jenn.lewis5789@gmail.com - Response times depend on support tier (see [LEGAL/SUPPORT_POLICY.md](#))

Documentation: - [docs/](#) — Complete documentation hierarchy - [docs/INDEX.md](#) — Documentation navigation - [docs/00_executive/GETTING_STARTED.md](#) — Quick start guide

9.2 Support Tiers

Standard Support (Included): - Email support - 72-hour response time - Documentation access

Business Support (Add-on): - 24-hour response time - Priority bug-fix handling - Pricing: Contact for quote

Enterprise Support (Add-on): - Same-business-day response - Scheduled support calls - Priority engineering resources - Pricing: \$5,000–\$20,000/month (see [LEGAL/SUBSCRIPTIONS.md](#))

Premium Support (Add-on): - White-glove service - Highest priority engineering - Flexible support scheduling - Pricing: Custom quote

See [LEGAL/SUPPORT_POLICY.md](#) for complete support tier details and [LEGAL/SUBSCRIPTIONS.md](#) for pricing.

10. Best Practices

10.1 Configuration Management

- **Version control** — Keep configurations in version control
- **Environment-specific** — Use separate configs for development, staging, production
- **Secrets management** — Never commit secrets; use environment variables or secure vaults

10.2 Testing

- **Test in development** — Always test updates in a development environment first
- **Backup before updates** — Backup configurations and models before major updates
- **Version pinning** — Pin dependency versions for reproducible builds

10.3 Monitoring

- **Logging** — Review logs regularly for errors or warnings
 - **Performance monitoring** — Monitor model performance and pipeline execution times
 - **Resource usage** — Monitor CPU, memory, and GPU usage
-

11. Troubleshooting

11.1 Common Issues

Installation Issues: - Check Python version (3.9+ required) - Verify all dependencies are installed
- Check system requirements

Configuration Issues: - Review configuration file syntax (YAML) - Check file paths and permissions
- Verify data source connectivity

Performance Issues: - Check system resources (CPU, RAM, GPU) - Review model configuration parameters
- Consider data preprocessing optimizations

11.2 Getting Help

Support Process: 1. **Check documentation** — Review relevant documentation first 2. **Gather information** — Collect logs, configs, and error messages 3. **Contact support** — Email support with detailed information 4. **Provide context** — Include version, environment, and steps to reproduce

12. Next Steps

12.1 Recommended Reading

1. **Getting Started** — [docs/00_executive/GETTING_STARTED.md](#)
2. **Architecture Overview** — [docs/ARCHITECTURE.md](#)
3. **Configuration Basics** — [docs/01_tutorials/configuration/CONFIG_BASICS.md](#)

4. First Pipeline Run — [docs/01_tutorials/pipelines/FIRST_PIPELINE_RUN.md](#)

12.2 Advanced Topics

- Feature Engineering — [docs/01_tutorials/pipelines/FEATURE_ENGINEERING_TUTORIAL.md](#)
 - Model Training — [docs/01_tutorials/training/BASIC_TRAINING.md](#)
 - Trading Integration — [docs/01_tutorials/trading/ALPACA_INTEGRATION.md](#)
-

Contact

For onboarding assistance or questions:

Jennifer Lewis

Fox ML Infrastructure LLC

Email: jenn.lewis5789@gmail.com

Subject: *Onboarding Inquiry — Fox ML Infrastructure*

Related Documents

- [LEGAL/ENTERPRISE_DELIVERY.md](#) — Repository structure and delivery model
- [LEGAL/SUPPORT_POLICY.md](#) — Support tiers and response times
- [LEGAL/RELEASE_POLICY.md](#) — Versioning and update policies
- [LEGAL/SECURITY.md](#) — Security practices and data handling
- [docs/INDEX.md](#) — Complete documentation navigation